

Linux perf

Locate performance bottleneck with low overhead

Geoffrey Papaux,
May 18, 2017

geoffrey.papaux@deltaww.com

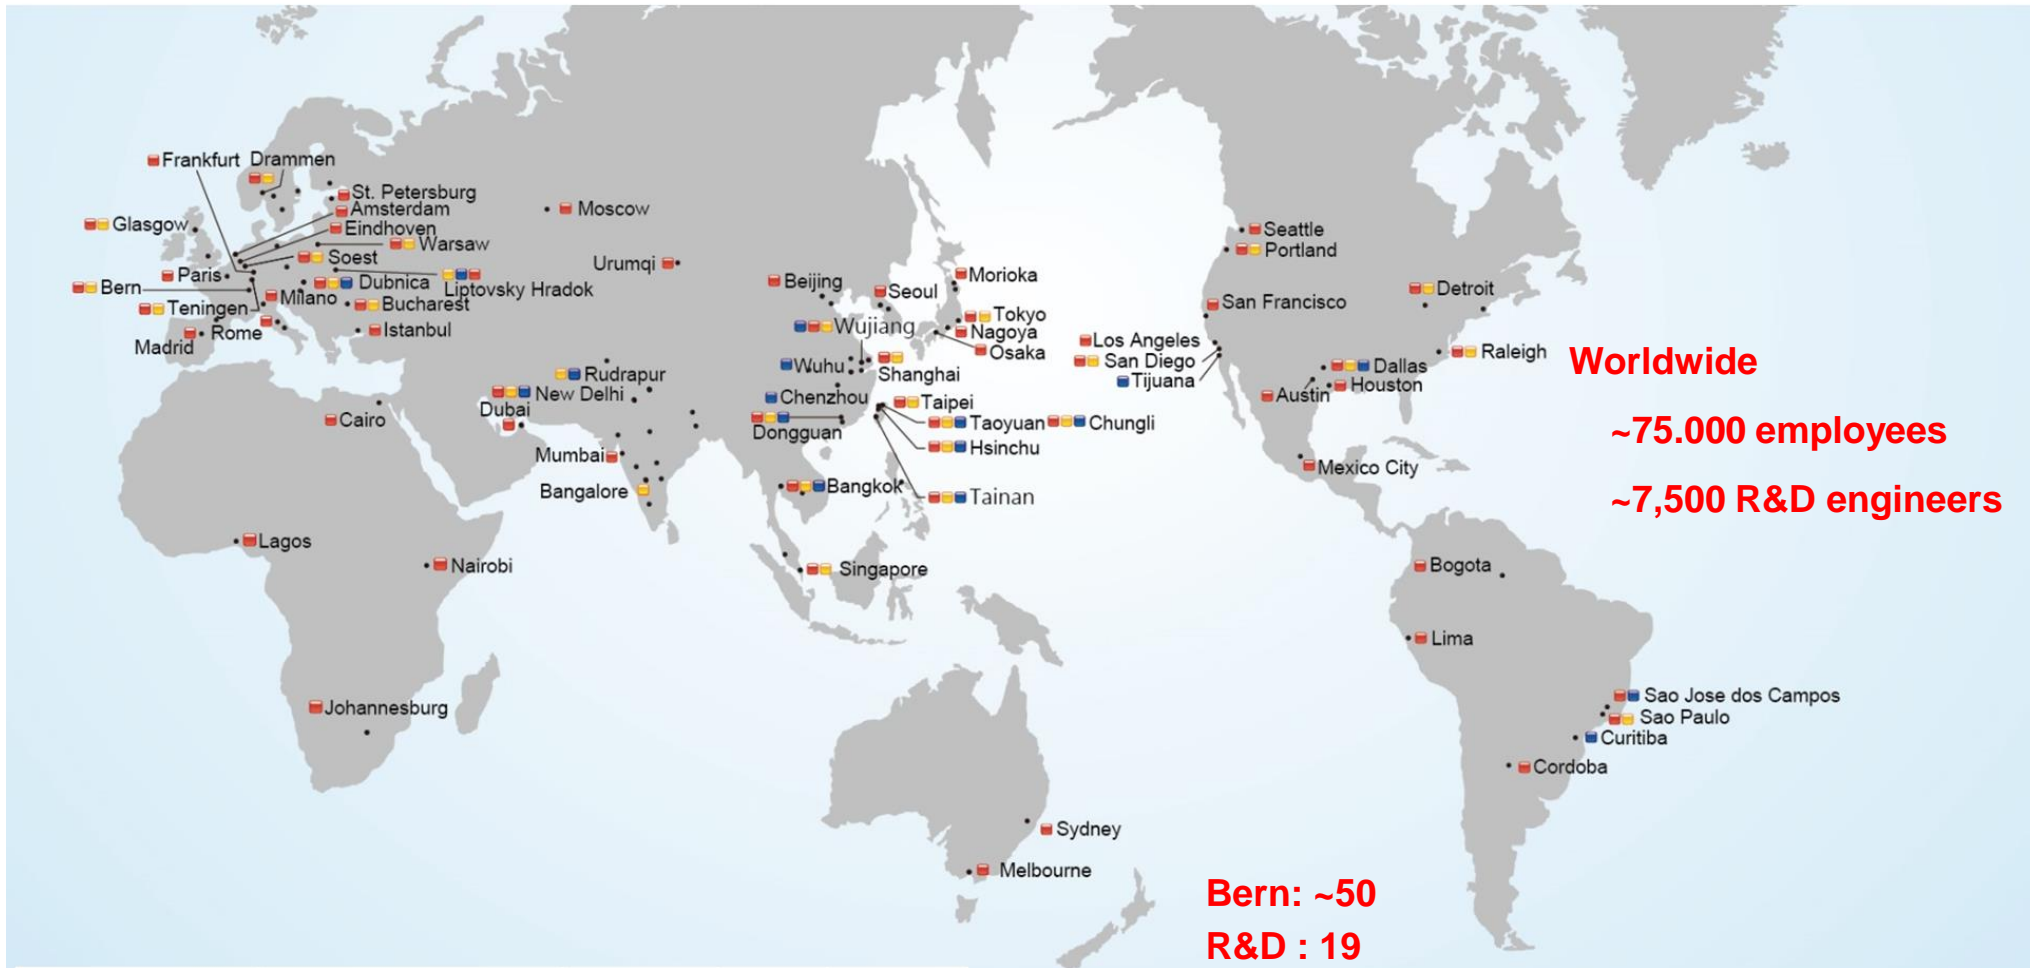


- About
- Profiling
- Performance Counters
- Perf
 - Architecture
 - Usage
 - Advanced options & use
 - Typical issues
- Conclusion

- Bachelor / Master @ HEIA-FR / HES-SO
- Embedded SW Engineer @ Delta Energy Systems
- «Experience» with perf ?
 - Bachelor Thesis: *Perf-based Profiles for GCC*
 - Master Thesis: *Virtualization on ARM*
 - Occasionally at Delta 😊



Delta Worldwide



	Asia-Pacific (China)	Americas	EMEA	Total
■ Sales Offices	96 (52)	20	37	153
■ Plant Sites	34 (22)	4	2	40
■ R&D Centers	43 (23)	7	11	61

Software Engineer: 9
Hardware Engineer : 2
Test Engineer: 2
Support : 3

Business Categories



Power Electronics

- Embedded Power Supplies
- Mobile Power Supplies
- Industrial and Medical Power Supplies
- Fans and Thermal Management
- Electronic Components for ICT Equipment



Energy Management

- Industrial Automation
- **Telecom Power Systems** **Delta Bern**
- UPS & Datacenter Infrastructure
- Automotive Electronics & EV Charging
- Renewable Energy
- Energy Storage Systems



Smart Green Life

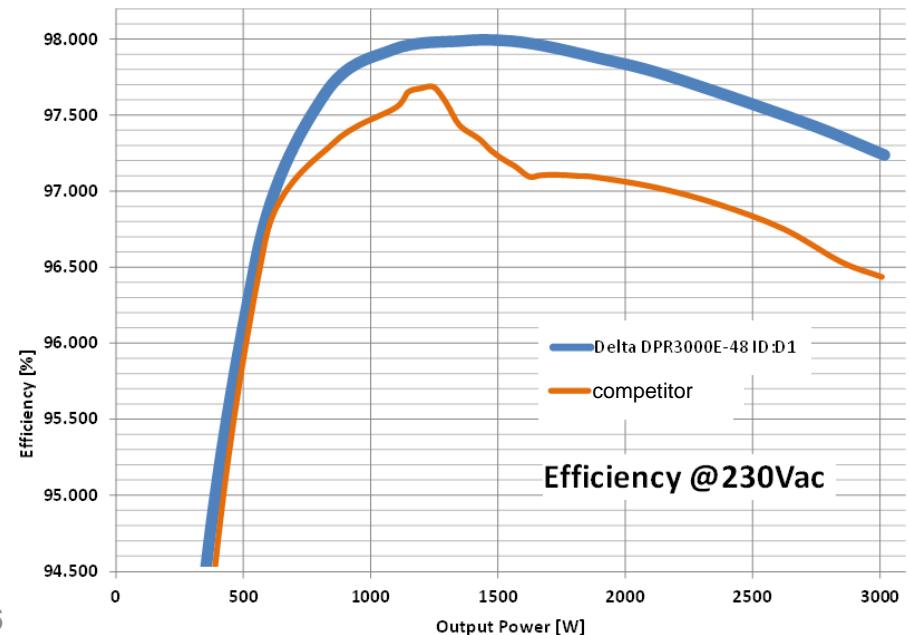
- Networking Systems
- Display & Visualization
- LED Lighting
- Healthcare Devices
- **Innervie**
- **vivitek**



- 4th Gen Power System Controller ORION
- Color 2" Touch screen
- ~100'000 units per year
- Software continuously enhanced
 - > 40 software version
 - > 700'000 lines of code
 - > 2000 configuration parameters
 - > 200 web pages
- > 50 different power converter or extension modules



- 48V DC Rectifier 3000W
- **98% efficiency**
- 56.8 W/inch³ power density



Profiling ?

When and why do I need it ?

HELP! My app is slow

- Slower than before
 - after adding a new functionality
 - after upgrading an external library
 - with a special kind of input / workload
 - after changing the hardware platform
- Needs to run faster
 - to satisfy timing constraints
 - to use less resources (save cloud costs)
- **=> I want to locate the bottleneck!**

- Art of collecting and analysing data
 - *dynamic* analysis (vs *static* analysis)
- Locate **hot spots**
 - where is the execution time spent ?
 - how often is this function called and from where ?
- Resolution
 - function level
 - instruction level
- Two categories: instrumentation or sampling

- The monitored program is modified
 - at source code level (`clock_gettime()`)
 - at compile time (`gprof`)
 - at execution time (`valgrind`)

(+) provides exact measurements

(-) high overhead ($\geq 30\%$ for `gprof`)

(-) need special binary / compilation (not for `valgrind`)

(-) program behavior is altered

Or **statistical profiling**

- The program runs unmodified
- A tool records the program state at regular intervals
 - by probing call stack
 - by using hardware capabilities

(+) very low overhead (1-2%)

(+) no special binaries needed

(+) behavior is unaltered

(-) not 100% accurate (but usually enough)

Profiling is not straightforward

- Today's architectures are complex
- Bottleneck is not only about instructions executed
 - memory access ? cache L1, L2, L3 ?
 - branch prediction ?
 - Instruction pipeline ?
- And not only hardware
 - context switch ?
 - CPU migration ?

Hardware to the rescue

- CPU registers counting hardware events
 - CPU Cycles
 - Branch-misses / Cache-misses
 - and many more
- Supported by Intel, AMD, ARM, PowerPC, ...
- Architecture specific
- Often referenced as
 - *PMU (Performance Monitoring Unit)*
 - *PMC (Performance Monitoring Counters)*

THE performance analysis tool for Linux.

*perf provides rich generalized abstractions
over hardware specific capabilities*

https://perf.wiki.kernel.org/index.php/Main_Page

- **Hardware** counters from PMU, but not only.
- Lot of **software** counters:
 - Context-switches
 - CPU Migrations
 - Kernel tracepoints
 - Page faults
 - ...

- perf is a *framework*
- **Back-end** must be enabled in your kernel config:
 - > General setup
 - [*] Kernel Performance Events And Counters
- **Front-end** is in Linux Kernel source tree
 - In directory tools/perf
- Wait. User-space application in Kernel Source ?
 - => tightly coupled to Linux kernel ABI
 - => require exact version matching

- From your distribution

```
apt-get install linux-tools-common linux-tools-generic
```

```
pacman -S perf
```

```
dnf install perf
```

- As a Yocto package

```
IMAGE_INSTALL = "perf"
```

- Compile from source

```
make -C tools/perf
```


Using perf

- git-like interface
- Usage via subcommands
`perf <subcommand>`
- List all subcommands
`perf`
- Getting help about a subcommand
`perf help <subcommand>`

Command	Description
list	list of available events
stat	collect events while running a command
record	record events in perf.data file
report	analyse a perf.data file
script	scripting interface for processing perf.data file
top	performance counters in real-time
bench	micro-benchmarking framework
...	...

- 25+ commands

- List available commands

```
$ perf
```

```
usage: perf [--version] [--help] [OPTIONS] COMMAND [ARGS]
```

The most commonly used perf commands are:

annotate	Read perf.data (created by perf record) and display annotated code
archive	Create archive with object files with build-ids found in perf.data
bench	General framework for benchmark suites
buildid-cache	Manage build-id cache.
buildid-list	List the buildids in a perf.data file
config	Get and set variables in a configuration file.
data	Data file related processing
diff	Read perf.data files and display the differential profile
evlist	List the event names in a perf.data file
ftrace	simple wrapper for kernel's ftrace functionality

[...]

- List available events (hardware and software)

```
$ perf list
```

List of pre-defined events (to be used in -e):

branch-instructions OR branches	[Hardware event]
branch-misses	[Hardware event]
cache-misses	[Hardware event]
cache-references	[Hardware event]
cpu-cycles OR cycles	[Hardware event]
instructions	[Hardware event]
alignment-faults	[Software event]
context-switches OR cs	[Software event]
[...]	

- Collect counter statistics while running a command

```
$ perf stat <command>
```

```
Performance counter stats for '<command>':
```

7934.312649	task-clock:u (msec)	#	1.000 CPUs utilized
0	context-switches:u	#	0.000 K/sec
0	cpu-migrations:u	#	0.000 K/sec
106	page-faults:u	#	0.013 K/sec
35,688,627,718	cycles:u	#	4.498 GHz
92,920,415,499	instructions:u	#	2.60 insn per cycle
14,406,550,955	branches:u	#	1815.728 M/sec
258,189,358	branch-misses:u	#	1.79% of all branches

```
7.934377499 seconds time elapsed
```

Try again with -d

- Display event counters in real time (top-like)

```
$ perf top
```

```
Samples: 482 of event 'cycles', Event count (approx.): 4563253774
```

Overhead	Shared Object	Symbol
43.84%	solver	[.] checkRow
16.96%	solver	[.] checkSquare
16.73%	solver	[.] checkColumn
8.50%	solver	[.] placeNum
2.57%	[kernel]	[k] format_decode
2.05%	solver	[.] goBack
1.81%	solver	[.] solveSudoku
0.51%	perf	[.] rb_next
[...]		

- Collect profiling data
- Data written in perf.data file
- Run an application and record

```
$ perf record <command>
```

```
[ perf record: Woken up 5 times to write data ]
```

```
[ perf record: Captured and wrote 1.244 MB perf.data (32133 samples) ]
```

- Collect system-wide

```
$ perf record -a
```

- Read and analyse perf.data previously recorded

```
$ perf report -n --stdio
```

```
# Samples: 32K of event 'cycles:u'
# Event count (approx.): 35983350419
```

Interactive analysis with
\$ perf report

```
#
# Overhead      Samples  Command  Shared Object  Symbol
# .....
#
```

40.69%	13073	solver	solver	[.] checkRow
24.49%	7869	solver	solver	[.] placeNum
17.71%	5689	solver	solver	[.] checkSquare
9.40%	3021	solver	solver	[.] checkColumn
4.16%	1337	solver	solver	[.] goBack
3.54%	1136	solver	solver	[.] solveSudoku

- Scope
 - -a system-wide collection (all CPUs)
 - -p <pid> from a running process
 - -t <tid> from a running thread
- Collect specific event(s)
 - -e <event>
- Collect at a given frequency (Hz)
 - -F <freq>
- Record call graph
 - -g

- «Cross-analysis»

- Collect on embedded system
- Analyse perf.data from another machine

```
$ perf report --objdump=<path> --symfs=<path-to-debug>
```

- Filter user / kernel

```
$ perf record -e cycles:u,cache-misses:k -a
```

- View report by source lines

```
$ perf report -s srcline
```

- Memory access analysis

```
$ perf mem [record|report]
```

- Built-in scripts

```
$ perf script -l
```

- Generate script

```
$ perf script -g python
```

```
$ python perf-script.py
```

- Example: record failed syscalls

```
$ perf script record failed-syscalls
```

```
$ perf script report failed-syscalls
```

- Dynamic tracing

```
$ perf probe /lib/libc.so.6 malloc
```

```
$ perf record -g -e probe_libc:malloc -aR
```

```
$ perf report
```


- Scheduler

```
$ perf sched record
```

```
$ perf report
```

- «Timechart»:

```
$ perf timechart
```




The Performance Monitors provide approximately accurate count information. To keep the implementation and validation cost low, a reasonable degree of inaccuracy in the counts is acceptable. ARM does not define a reasonable degree of inaccuracy.

ARM Architecture Reference Manual (ARMv7-A)

Typical issues

- Stack not working (=> `-fno-omit-frame-pointer`)
- Debug symbols not available
- No PMU support (typical `<not supported>` message)
 - virtual machine / cloud
 - not available on your hardware
 - not implemented for your target
 - sometimes just not enabled in device tree
- Sudo is your friend, even for perf list

Conclusion



Conclusion

- perf is a powerful, low overhead profiling tool
- Actively developed <https://kernelnewbies.org/LinuxChanges>
 - Getting new functionality with every kernel release
 - Dedicated section in change log
 - Listed in *Prominent Features* in 4.11, 4.10, 4.7, 4.4
- Available for Linux > 2.6.31
- ! Understand what you are measuring
- Make sure it is working before making any conclusion

- Perf wiki: <https://perf.wiki.kernel.org>
- Kernel documentation: tools/perf/Documentation/
- Architecture reference manuals
 - INTEL: <http://www.intel.com/Assets/PDF/manual/253669.pdf>
 - ARM: http://support.amd.com/us/Processor_TechDocs/31116.pdf
- Brendan Gregg, Netflix Performance Engineer
<http://www.brendangregg.com/perf.html>

Thanks for your attention



Smarter. Greener. Together.

To learn more about Delta, please visit www.deltaww.com.

